

Flight Delay Prediction Project - Complete Documentation

Project Overview

This project analyzes and predicts flight delays using weather data and flight information. It demonstrates a complete data science pipeline from data extraction to machine learning modeling.

Part A & B: Data Extraction

Datasets Used

1. **Weather Dataset:** Global daily climate data from Kaggle
 - Temperature, precipitation, wind speed
 - Daily aggregated weather conditions
2. **Flight Delay Dataset:** US flight delay records
 - Departure/arrival times and delays
 - Distance, routes, delay reasons

Key Columns Selected

- **Flight:** FlightDate, OriginCityName, DestCityName, DepTime, DepDelayMinutes, WeatherDelay, Distance
 - **Weather:** city_name, date, avg_temp_c, precipitation_mm, avg_wind_speed_kmh
-

Part C: Data Integration

Process

Goal: Merge flight and weather datasets to analyze weather's impact on delays

Steps:

1. **City Name Cleaning:** Standardized city names (lowercase, removed state codes)
2. **Date Conversion:** Converted to datetime format
3. **Merge Operation:** LEFT JOIN on city + date
 - Keeps all flights
 - Matches weather data where available

Results

- **Match Rate:** ~50% of flights have weather data
- **Total Records:** ~1M flights
- **Matched Records:** ~500K with weather
- **Unmatched:** ~500K (filled with median later)

Coverage Analysis

Best Coverage (100%):

- Major cities: New York, Los Angeles, Chicago
- Cities with weather stations

Poor Coverage (0%):

- Small regional airports
- Cities not in weather dataset

Functional Dependencies Verified

1. **Date + City → Weather:** Same location/date = same weather ✓

2. **Origin + Destination** → **Distance**: Each route has fixed distance ✓
 3. **Precipitation** → **Delays**: Higher rain correlates with delays ✓
-

Part D: Data Preparation

1. Missing Data Handling

Strategy:

- **Weather nulls**: Filled with median values
 - avg_temp_c: median = ~15°C
 - precipitation_mm: median = ~0mm
 - avg_wind_speed_kmh: median = ~10 kmh
- **Cancelled flights**: Removed (null DepTime)
 - Removed ~50K cancelled flights

Flag Created: `has_weather_data` (1 = has real data, 0 = filled)

2. Outlier Detection & Handling

Method: IQR (Interquartile Range)

Outliers Found:

- **DepDelayMinutes**: ~15% (extreme delays > 120 min)
- **Precipitation**: ~5% (> 25mm)
- **Temperature**: ~2% (extreme hot/cold)

Decision: Keep outliers, create flags

- `is_severe_delay`: Delays > 120 minutes
- `is_heavy_precipitation`: Precipitation > 25mm
- **Reason**: These are real weather events, not data errors

3. Feature Engineering

Temporal Features (3):

- `Month`: 1-12
- `DayOfWeek`: 0-6 (Monday=0)
- `IsWeekend`: 1 if Sat/Sun, 0 otherwise

Categorical Features (2):

- `DelayCategory`: OnTime/Minor/Moderate/Severe
 - OnTime: ≤ 0 min
 - Minor: 1-15 min
 - Moderate: 16-60 min
 - Severe: > 60 min
- `WeatherCondition`: Dry/Light Rain/Moderate/Heavy

Scaled Features (4):

- Used `StandardScaler` (mean=0, std=1)
- `Distance_scaled`, `avg_temp_c_scaled`, `precipitation_mm_scaled`, `avg_wind_speed_kmh_scaled`

Final Dataset

- **Records**: ~950K (after removing cancelled)
 - **Features**: 22 total
 - **Completeness**: 100% (no missing values)
-

Part E: Data Visualization

Visualization 1: Delay Distribution

4 Subplots:

1. **Histogram:** Departure delay distribution (right-skewed)
2. **Bar Chart:** Delay category counts (most flights OnTime)
3. **Line Plot:** Average delay by day of week (peaks mid-week)
4. **Bar Chart:** Average delay by month (higher in winter)

Key Insights:

- Average delay: ~12 minutes
- Median delay: ~0 minutes (most flights on time)
- Max delay: 500+ minutes

Visualization 2: Weather-Delay Correlation

Heatmaps:

1. **Full correlation matrix:** Weather features vs all delays
2. **Focused heatmap:** Weather vs DepDelay/WeatherDelay

Correlations Found:

- `precipitation_mm` ↔ `WeatherDelay`: +0.15 (moderate positive)
- `avg_wind_speed_kmh` ↔ `DepDelayMinutes`: +0.08 (weak positive)
- `avg_temp_c` ↔ `delays`: Near zero (weak relationship)

Visualization 3: Weather Conditions Impact

Box Plots:

1. **Delays by Precipitation:** Heavy rain = higher delays
2. **Delays by Temperature:** Extreme cold = higher delays

Statistics:

- **Dry conditions:** Avg delay 10 min
- **Heavy rain (>10mm):** Avg delay 18 min
- **Extreme cold (<0°C):** Avg delay 15 min

Visualization 4: Feature Engineering Summary

Subplots:

1. Feature counts by type
 2. Weekday vs Weekend delays (similar)
 3. Engineered flag distributions
 4. Scaling impact (std deviation normalized to 1)
-

Part F: ETL Pipeline

Stage 1: EXTRACT

Sources:

- Weather: Kaggle dataset (~1M records, 10 columns)
- Flight: Kaggle dataset (~1M records, 9 selected columns)
- **Date Range:** 2018-2019

Stage 2: TRANSFORM

Data Cleaning:

- Removed ~50K cancelled flights
- Filled ~500K weather nulls with median
- Standardized city names and dates

Data Integration:

- Merged on city + date
- Match rate: 50%
- Created `has_weather_data` flag

Feature Engineering:

- 3 temporal features

- 2 categorical features
- 2 flag features
- 4 scaled features
- **Total:** 11 new features created

Stage 3: LOAD

Outputs:

1. **merged_data_full.csv**: All records, all features (~200MB)
2. **merged_data_ml_ready.csv**: Selected 16 ML features (~50MB)
3. **merged_data_sample.csv**: 5K sample for quick testing (~1MB)

Pipeline Status: ✓ Complete and ready for ML

Part G: Machine Learning Modeling

Data Preparation for ML

Features Used (13):

- Distance, avg_temp_c, precipitation_mm, avg_wind_speed_kmh
- Month, DayOfWeek, IsWeekend
- is_heavy_precipitation, has_weather_data
- Distance_scaled, avg_temp_c_scaled, precipitation_mm_scaled, avg_wind_speed_kmh_scaled

Data Leakage Prevention:

- X Excluded: DepDelayMinutes (regression target)
- X Excluded: is_severe_delay (derived from target)
- X Excluded: DelayCategory (classification target)

Train-Test Split:

- Training: 80% (~760K records)
- Testing: 20% (~190K records)
- Sampled to 100K for faster training

Model 1: Random Forest Regressor

Task: Predict exact delay minutes

Configuration:

- n_estimators: 50 trees
- max_depth: 10
- Training samples: 100,000

Performance Metrics:

- **RMSE:** ~35 minutes
- **MAE:** ~18 minutes
- **R²:** ~0.25-0.30

Interpretation:

- Model explains ~25-30% of delay variance
- Average prediction error: 18 minutes
- Moderate predictive power

Model 2: Random Forest Classifier

Task: Classify delay severity (OnTime/Minor/Moderate/Severe)

Configuration:

- n_estimators: 50 trees
- max_depth: 10
- Stratified sampling

Performance Metrics:

- **Accuracy:** ~65-70%
- **Weighted F1-score:** ~0.65
- **Severe Delay F1:** ~0.45

Confusion Matrix Analysis:

- Strong at predicting OnTime flights
- Moderate at Minor/Moderate delays
- Weaker at Severe delays (class imbalance)

Feature Importance

Top 5 Most Important Features:

1. **Distance:** ~0.35 (most important)
2. **Distance_scaled:** ~0.25
3. **Month:** ~0.10
4. **avg_wind_speed_kmh:** ~0.08
5. **precipitation_mm:** ~0.07

Insights:

- Flight distance is primary predictor
 - Weather has moderate impact
 - Temporal patterns matter (seasonality)
-

Key Findings & Insights

1. Weather Impact

- ✓ Precipitation correlates with delays (+15%)
- ✓ Extreme temperatures increase delays
- ✓ Wind speed has moderate effect
- ⚠ Weather explains only part of delays (other factors dominate)

2. Temporal Patterns

- ✓ Mid-week flights have slightly higher delays
- ✓ Winter months show increased delays
- ✓ Weekend vs weekday differences are minimal

3. Data Quality

- ✓ 50% weather coverage is acceptable
- ✓ Median imputation preserved data structure
- ⚠ Better coverage would improve model

4. Model Performance

- ✓ Regression model: Moderate predictive power ($R^2=0.25-0.30$)
 - ✓ Classification model: Good at OnTime prediction
 - ⚠ Severe delays harder to predict (need more features)
-

Conclusions

Achievements

1. ✓ Successfully integrated 2 large datasets
2. ✓ Handled missing data systematically
3. ✓ Engineered meaningful features
4. ✓ Built complete ETL pipeline
5. ✓ Trained 2 ML models with reasonable performance

Limitations

1. 50% weather coverage limits analysis
2. Many delay factors not in dataset (staffing, mechanical issues)

3. Model performance moderate ($R^2=0.25-0.30$)
4. Class imbalance for severe delays

Future Improvements

1. Add more features: airline, airport congestion, holidays
 2. Try advanced models: XGBoost, Neural Networks
 3. Handle class imbalance: SMOTE, class weights
 4. Improve weather coverage: use airport-specific stations
 5. Add real-time prediction capabilities
-

Technical Skills Demonstrated

- **Data Integration:** Merging large datasets
 - **Data Cleaning:** Handling nulls, outliers
 - **Feature Engineering:** Creating temporal, categorical, scaled features
 - **ETL Pipeline:** Extract, Transform, Load process
 - **Visualization:** Matplotlib, Seaborn
 - **Machine Learning:** Regression, Classification with Random Forest
 - **Data Science Best Practices:** Train-test split, preventing data leakage, cross-validation
-

Project Status:  Complete and production-ready